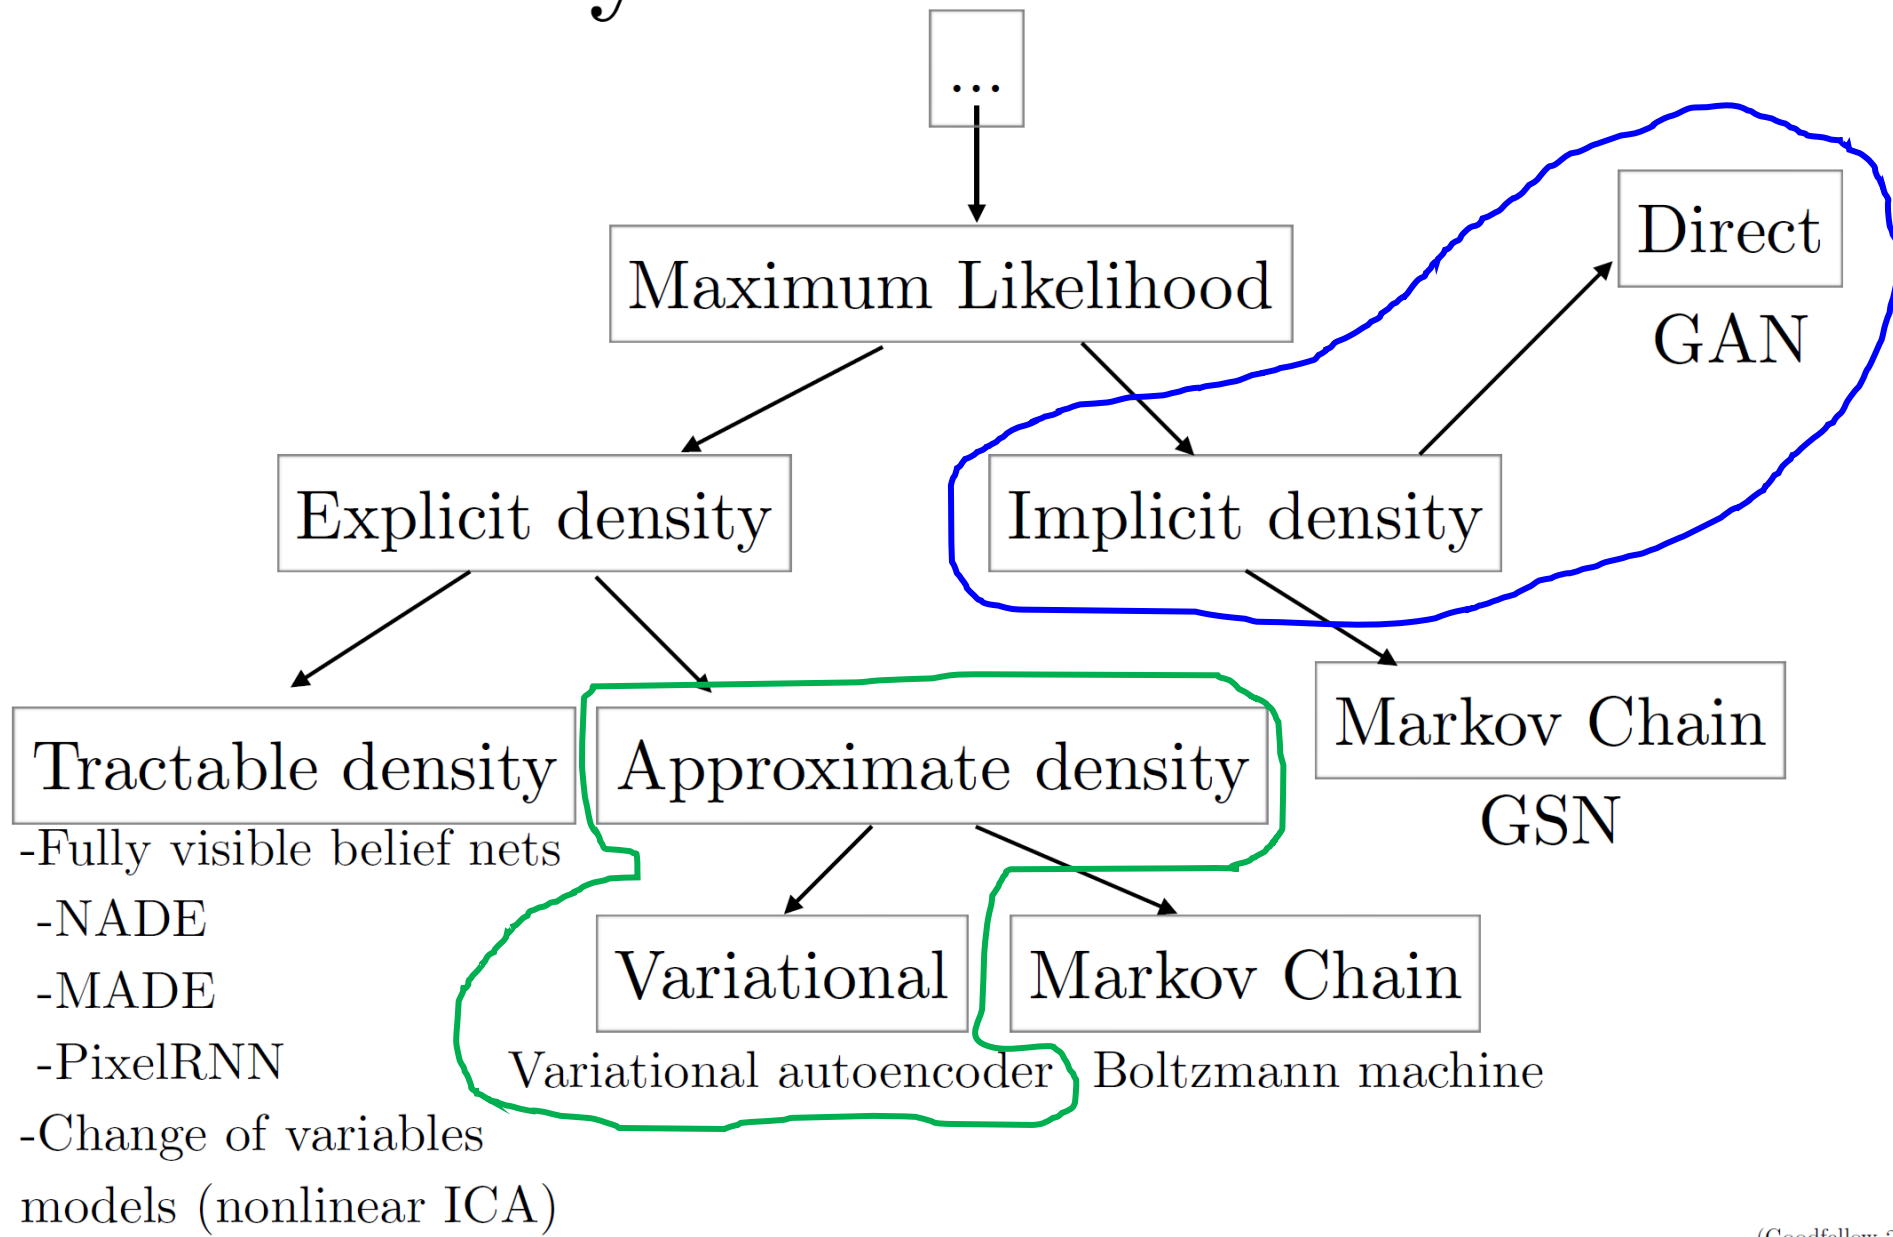


Variational Autoencoder

Generative Models

- **Generative Models**
- **Task:** Given a dataset of images $\{X_1, X_2, \dots\}$ can we learn the distribution of X ?
- Generative models are typically meant for modelling $P(X)$.
- Often applications seek for models which we can **sample** from.
 - Models that can generate random examples that follow the distribution of $P(X)$.

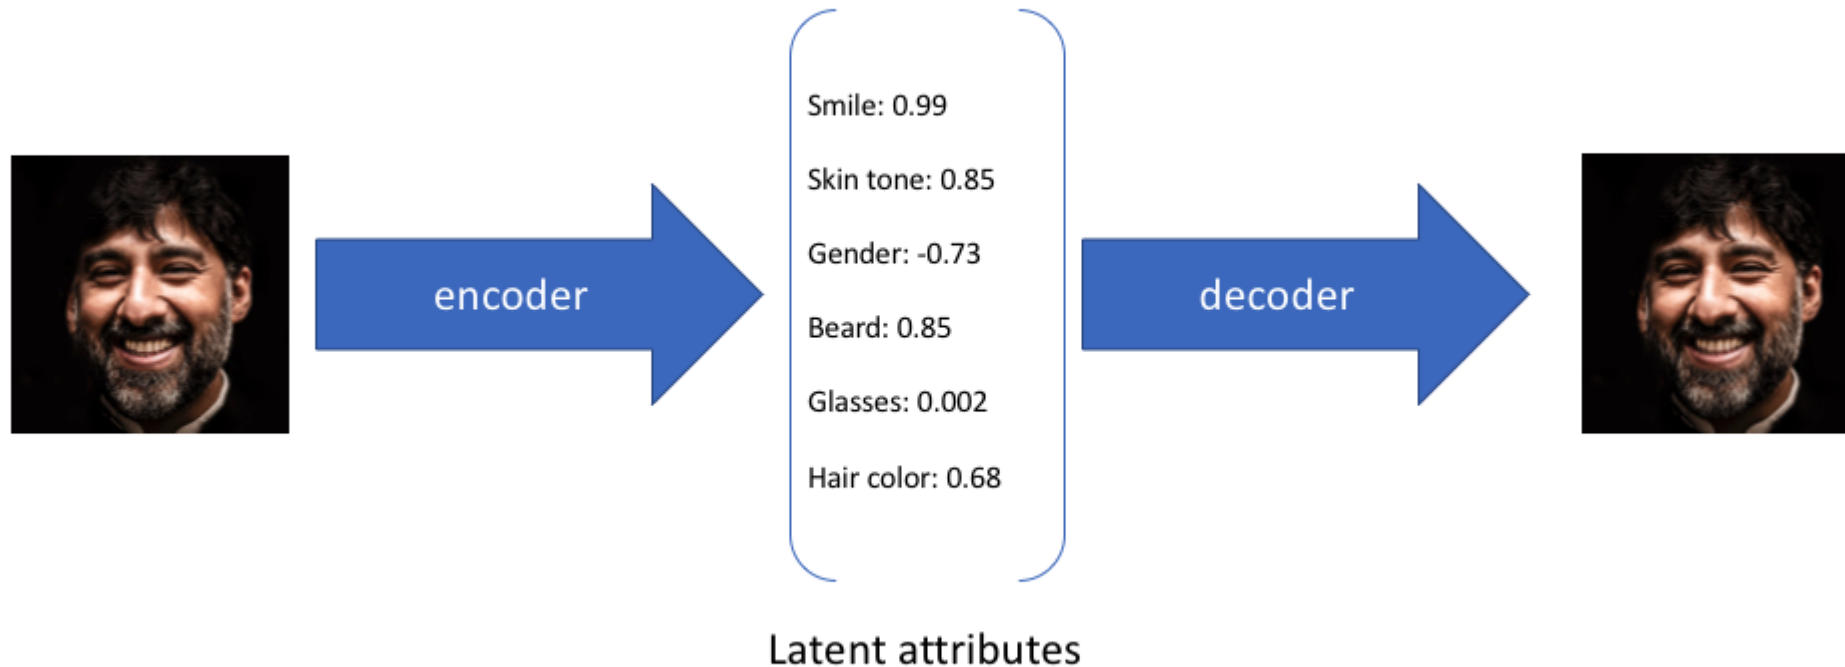
Taxonomy of Generative Models



Generative Models: GAN

- Remember that explicit form of $P(X|z)$ is not needed with input z in the latent space (features).
- But it is difficult to map back to the latent variable given a desired form of an output (an image).

Reconstruction from latent space variable is difficult



Variational Autoencoders





- Prior to GAN, variational autoencoders (VAEs) were meant for explicit Modelling of $P(X/z)$ with the given model parameters.
- $z \sim P(z)$, which can be sampled from Gaussian distribution.

$$P(X) = \int P(X|z; \theta) P(z) dz$$




Can we sample from a Gaussian Distribution to reconstruct the input ?



Smile (discrete value)



Smile (probability distribution)



VS.

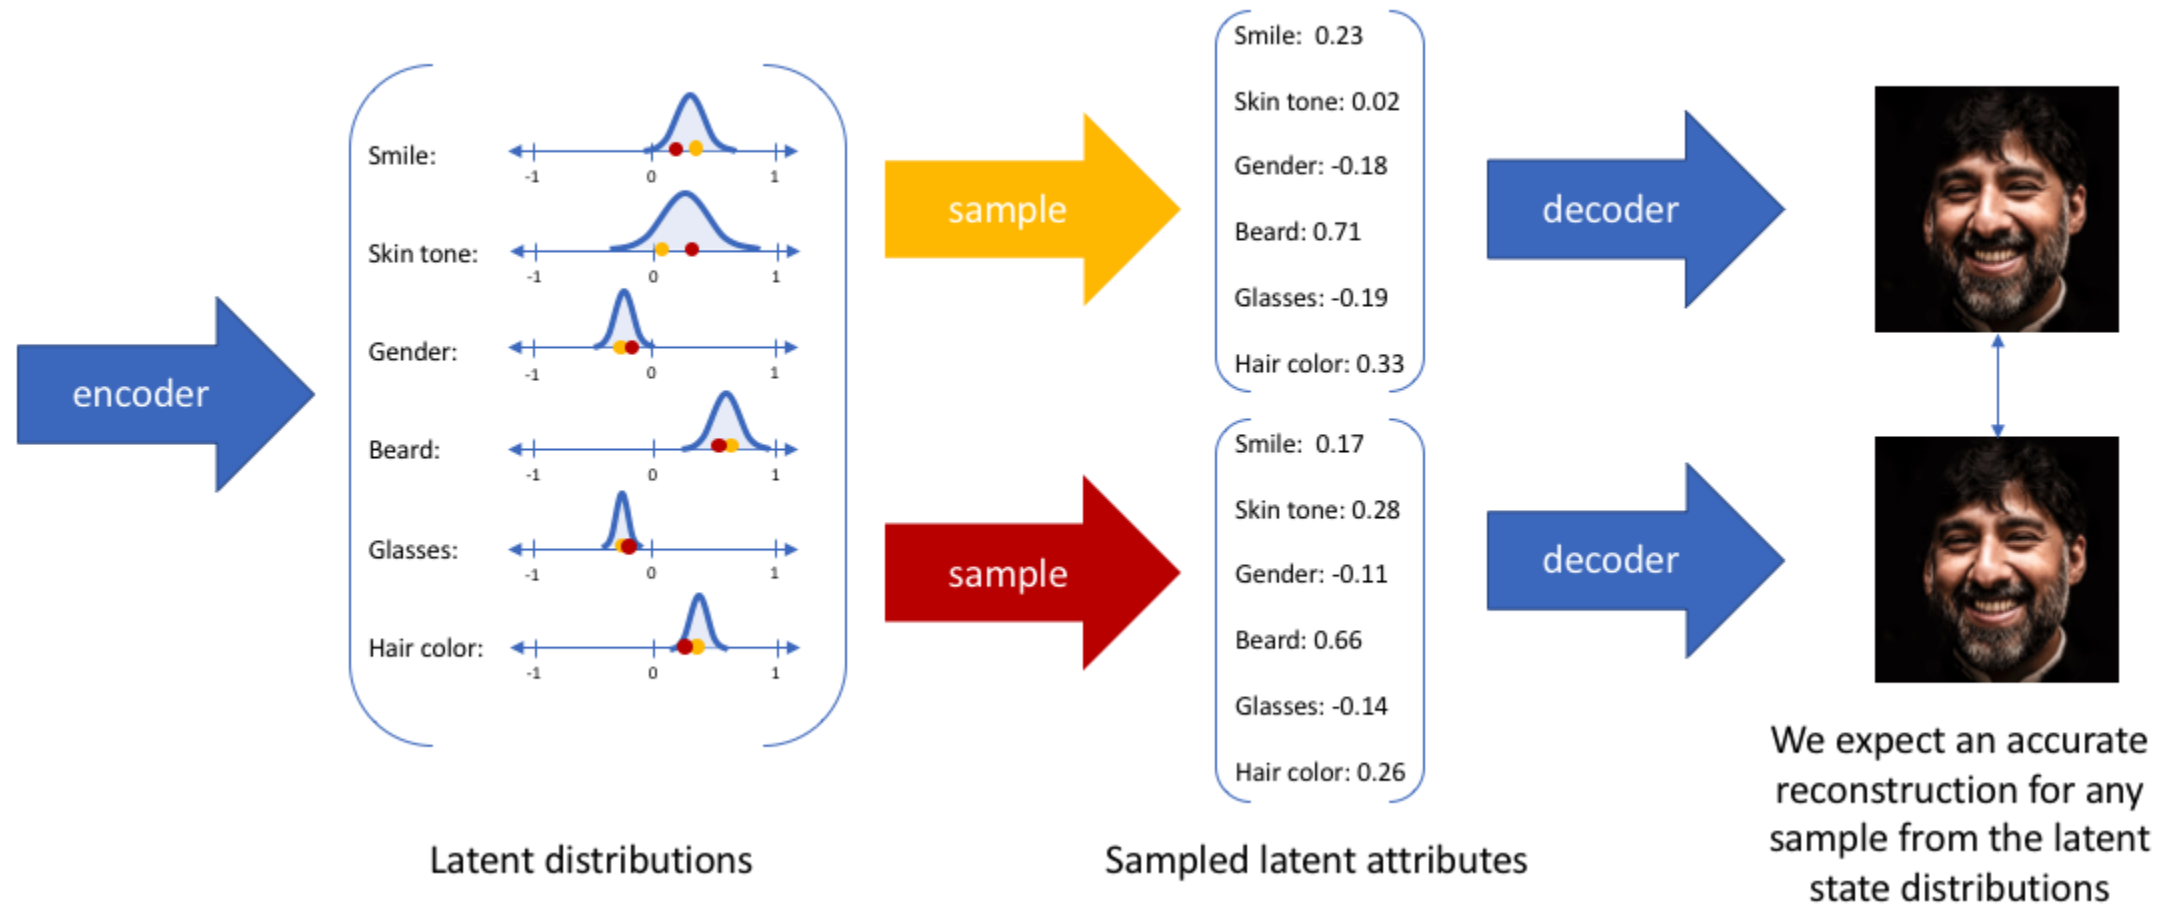
Variational Autoencoders

- Maximum Likelihood --- Find the model parameters to maximize $P(X)$, where X is the data.
- Approximate with samples of z

$$P(X) \approx \frac{1}{n} \sum_{i=0}^n P(X|z_i)$$

- So one requires large number of samples of z .
- For most of them $P(X|z) \approx 0$.
- Not practical computationally.

Approximate with samples of z : We want Accurate Reconstruction



Variational Autoencoders

- Is it possible to know which z will generate $P(X|z) \gg 0$?
- It amounts to learning a distribution $Q(z)$, where $z \sim Q(z)$ generates $P(X|z) \gg 0$.

Variational Autoencoders

- Suppose we can learn a distribution $Q(z)$, where $z \sim Q(z)$ generates $P(X|z) \gg 0$.
- We want to learn $P(X)$ such that $P(X) = E_{z \sim P(z)} P(X|z)$
 - not so practical. **Why ?**
- We will compute $E_{z \sim Q(z)} P(X|z)$ which is a more practical approach.

Variational Autoencoders

- How can we relate $E_{z \sim Q(z)} P(X|z)$ with $P(X)$?
- To know the relation, we first define Kullback–Leibler (KL) Divergence also known as relative entropy (measure).

$$D(Q(z) \| P(z|X)) = E_{z \sim Q(z)} [\log Q(z) - \log P(z|X)]$$

Relation Between $E_{z \sim Q(z)} P(X|z)$ and $P(X)$

- By Bayes rule

$$P(z|X) = \frac{P(X|z)P(z)}{P(X)}$$

- Apply logarithm on both sides

$$\log P(z|X) = \log P(X|z) + \log P(z) - \log P(X)$$

Relation Between $E_{z \sim Q(z)} P(X|z)$ and $P(X)$

- By definition of KL divergence

$$D(Q(z) \| P(z|X)) = E_{z \sim Q(z)} [\log Q(z) - \log P(z|X)]$$

- Substitute for $\log P(z|X)$ in the above expression.

$$\log P(z|X) = \log P(X|z) + \log P(z) - \log P(X)$$

Relation Between $E_{z \sim Q(z)} P(X|z)$ and $P(X)$

- KL divergence

$$D(Q(z) \| P(z|X)) = E_{z \sim Q(z)} [\log Q(z) - \log P(X|z) - \log P(z) + \log P(X)]$$

- By the properties of expectation function

$$E_{z \sim Q(z)} [Y + \text{constant}] = E_{z \sim Q(z)} [Y] + \text{constant}$$

- Therefore

$$D(Q(z) \| P(z|X)) = E_{z \sim Q(z)} [\log Q(z) - \log P(X|z) - \log P(z)] + \log P(X)$$

Relation Between $E_{z \sim Q(z)} P(X|z)$ and $P(X)$

$$D(Q(z) \parallel P(z|X)) = E_{z \sim Q(z)} [\log Q(z) - \log P(X|z) - \log P(z)] + \log P(X)$$

- Rearrange terms

$$\log P(X) - D(Q(z) \parallel P(z|X)) = E_{z \sim Q(z)} [\log P(X|z)] - E_{z \sim Q(z)} [\log Q(z) - \log P(z)]$$

$$\log P(X) - D(Q(z) \parallel P(z|X)) = E_{z \sim Q(z)} [\log P(X|z)] - D(Q(z) \parallel P(z))$$

Lower Bound on Log Probability

- Recall we want to maximize $P(X)$ with respect to model parameters.
- But we can not maximize $P(X)$ as we have no control \mechanism to maximize this probability from the latent variables.
- Now we have

$$\log P(X) - D(Q(z) \parallel P(z|X)) = E_{z \sim Q(z)}[\log P(X|z)] - D(Q(z) \parallel P(z))$$

- KL divergence D in the above expression is always > 0 . This means

$$\log P(X) > \log P(X) - D[Q(z) \parallel P(z|X)].$$

- So we maximize the lower bound on $\log P(X)$.

Lower Bound on Log Probability

- Hence the problem boils down to maximizing the following expression (lower bound on $\log P(X)$)

$$E_{z \sim Q(z)} [\log P(X|z)] - D(Q(z) || P(z))$$

Remember the assumption that

“we can learn a distribution $Q(z)$, where $z \sim Q(z)$ generates $P(X|z) \gg 0$.”

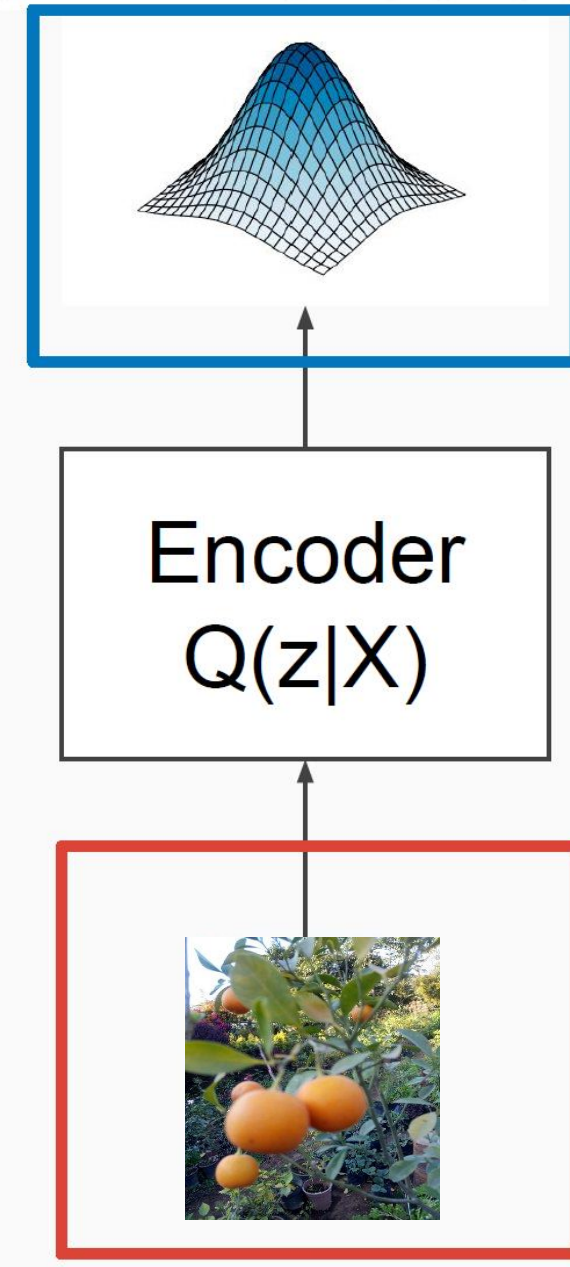
How do we get $Q(z)$?

How to get $Q(z)$?

- Model $Q(z|X)$ with a neural network.
- Assume $Q(z|X)$ to be a Gaussian

$$N(\mu, c \cdot I)$$

- Neural network **outputs** the **mean** μ , and a diagonal covariance matrix $c \cdot I$.
- **Input:** Image
- **Output:** Distribution: Two vectors μ and c .
- Call $Q(z|X)$ the '**Encoder**'.



Variational Autoencoder – Loss Function

$$\log P(X) - D(Q(z) \parallel P(z|X)) = E_{z \sim Q(z)} [\log P(X|z)] - D(Q(z) \parallel P(z))$$

- Let us convert the lower bound to a loss function:
- Model $P(X|z)$ with a neural network.
- Let $f(z)$ be the network output.
- Assume X to be independent identically distributed random variable in a Gaussian distribution.
 - $X = f(z) + \eta$, where $\eta \sim N(0, I)$.
- Then the problem is reduced to minimizing the error of regression

$$\|X - f(z)\|^2$$

- Call $P(X|z)$ the Decoder.


Variational Autoencoder – Loss Function ...

- If $P(z) \sim N(0,1)$ then $D[Q(z|X) || P(z)]$ has a closed form solution.
- So the loss function can be viewed as

$$L = \|X - f(z)\|^2 - \alpha D[Q(z|X) || P(z)]$$


Training

- Note that the model is stochastically generating output which means that even for the same input, mean and standard deviations , the actual encoding may vary slightly on every single pass.
- This is due to random sampling from the encoder's output.




Training

- If only $\|X - f(z)\|^2$ is chosen for optimization on a dataset, it results in creation of distinct clusters for different image classes.
- Therefore, the decoder is not able to produce any variance from an input image, it learns to replicate the input, as in standard autoencoder.




Training

- If only KL loss $D[Q(z|X) || P(z)]$ is chosen for optimization, it results in encoded vectors that are densely placed near the center of the latent space
- The decoder is not able to infer properly due to denseness in the latent space with not much variation.



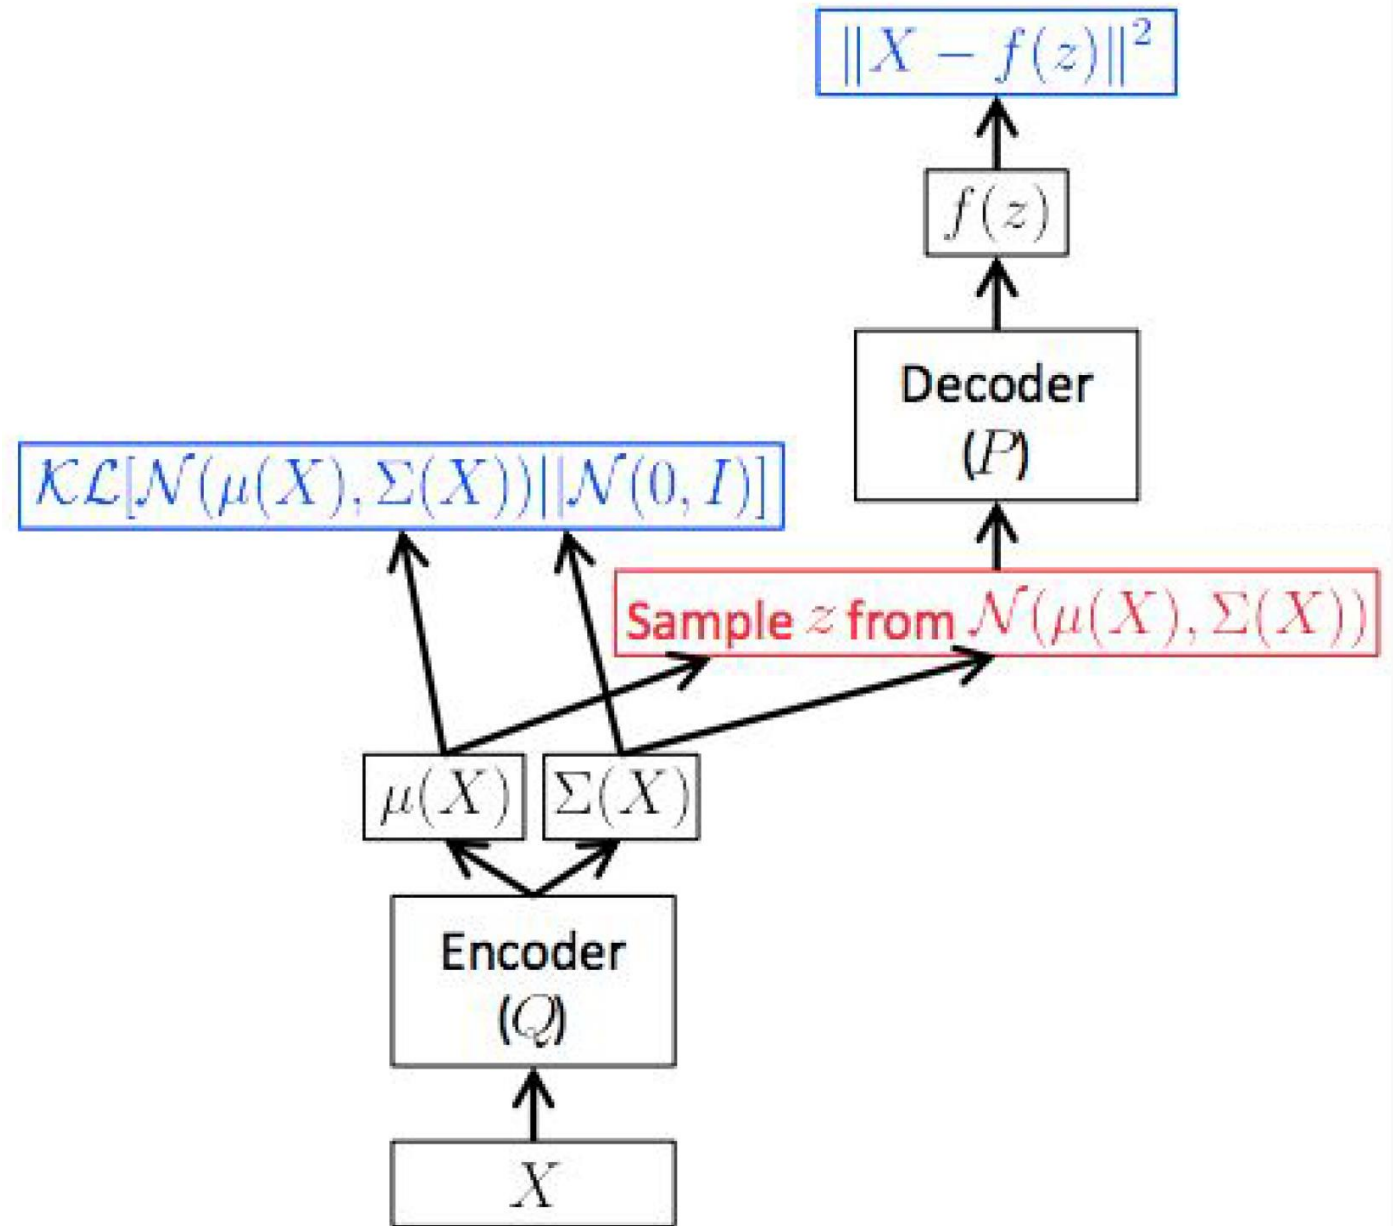
Training

- Optimizing the two together helps in generation of a latent space which maintains the similarity of nearby encoded vectors by clustering them together.
- And is very dense near the latent space origin.
- So the decoder is able to learn enough about variances in encoded vectors.



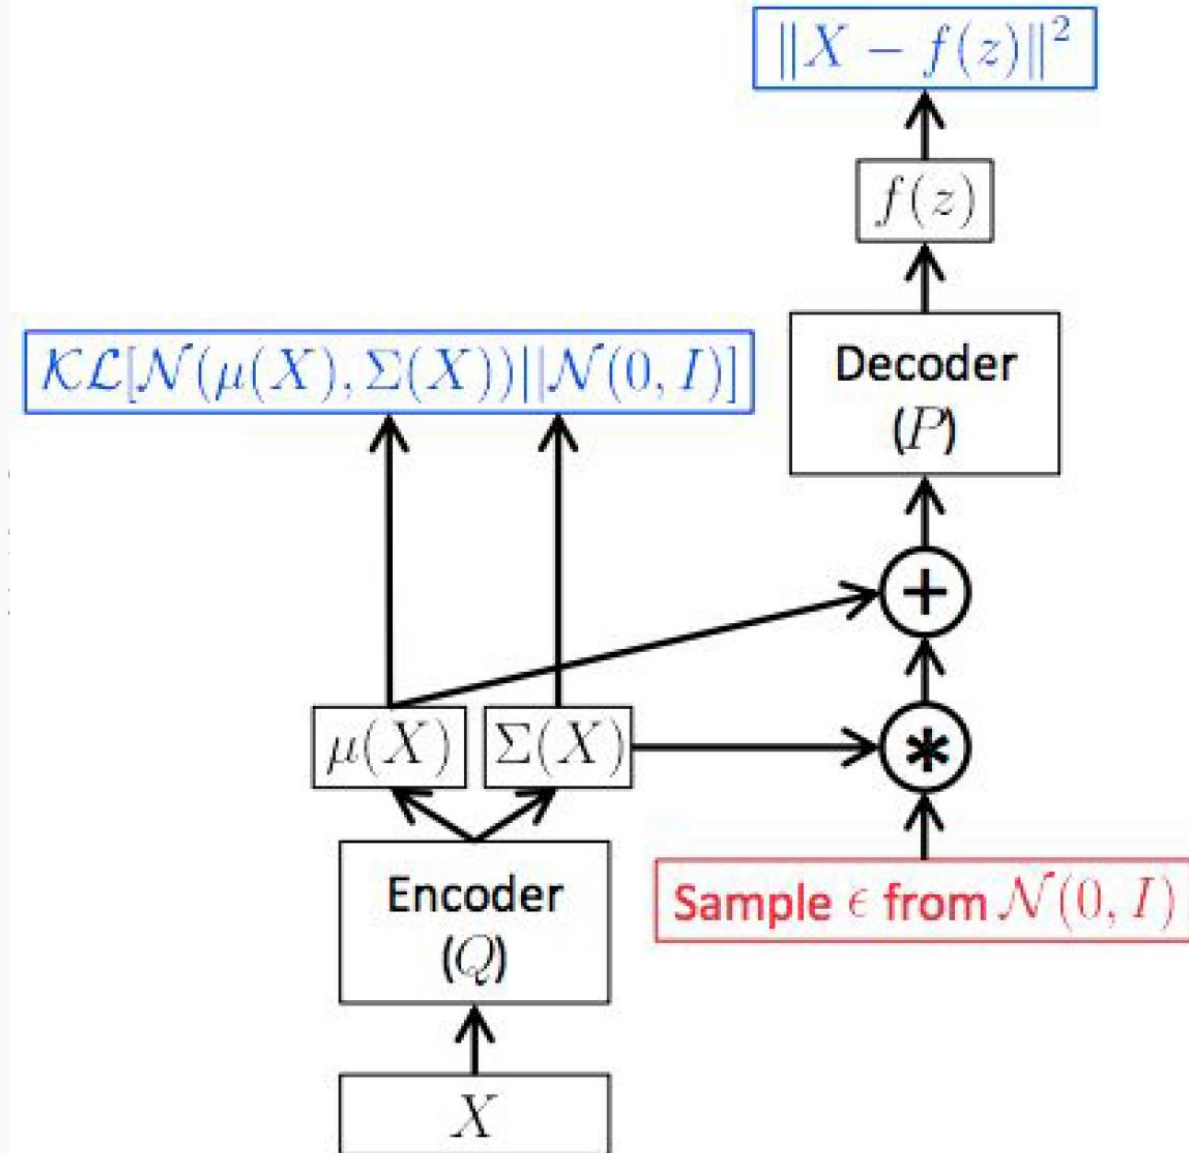
Training

- Training the Decoder is easy using the standard backpropagation.
- Training of encoder requires a better thinking. How to get the distribution ?



Reparameterization

- How to effectively backpropagate through the z samples to the Encoder?
- **Reparameterization Trick**
- $z \sim \mathcal{N}(\mu, \sigma)$ is equivalent to $\mu + \sigma \cdot \epsilon$, where $\epsilon \sim \mathcal{N}(0, 1)$
- Once this is done, we can easily backpropagate the loss to the Encoder.



Training


- Given a dataset of examples $X = \{X_1, X_2, \dots\}$
- Initialize parameters for Encoder and Decoder
- **Repeat till convergence:**
- Take a random minibatch X_M of M examples from X
- $\epsilon \leftarrow$ Sample M noise vectors from $N(0, I)$
- Compute the loss $L(X_M, \epsilon, \theta)$ after a forward pass in the neural network.
- Use gradient descent on L to update Encoder and Decoder.

Testing

- To evaluate the performance of VAE on generating a new sample.
- Sample $z \sim N(0, I)$ and pass it through the Decoder.
- No role of encoder as the latent variable itself is passed through the decoder.
- No good measure, relies on visual inspection.

VAE on MNIST Dataset

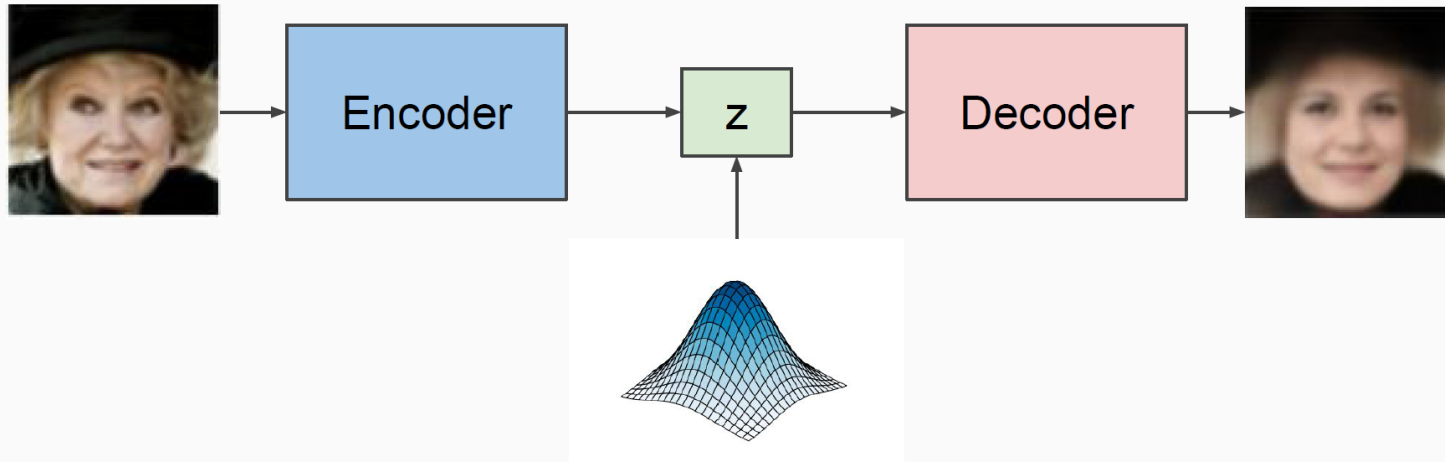
- As you see, distinct digits smoothly transform from one digit to another.
- This smooth transformation is useful when we want to interpolate between two observations, like a smiling face and a laughing face, a face without and with spectacles.



A Comprehensive Repository of Generative Model Codes


- <https://github.com/wiseodd/generative-models>

VAE and GAN: Comparison




VAE

GAN



VAE and GAN: Comparison




GAN

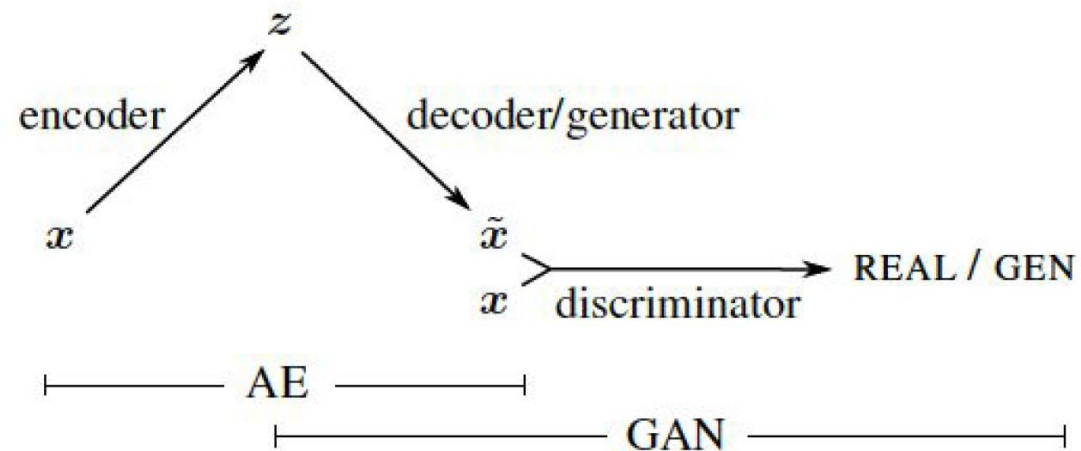
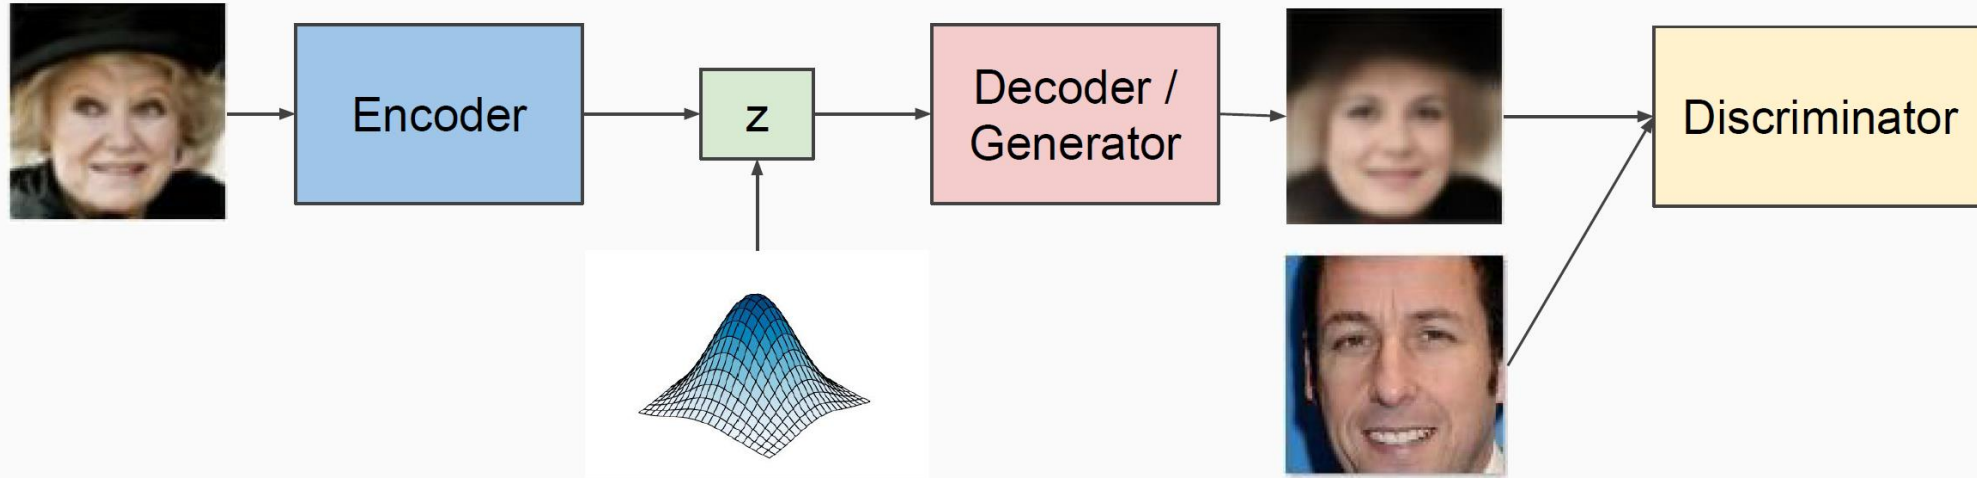
✓ : Very sharp images

X : Given an X **difficult** to find z . (Need to backprop.)

✓ / X : No explicit $P(X)$.





Combined VAE + GAN



KL Divergence L_2 Difference

$$\mathcal{L} = \mathcal{L}_{\text{prior}} + \mathcal{L}_{\text{llike}}^{\text{Disl}} + \mathcal{L}_{\text{GAN}}$$

Results



VAE_{Dis_l}: Train a GAN first, then use the discriminator of GAN to train a VAE.

VAE/GAN: GAN and VAE trained together.

Acknowledgement

- Sincere thanks to Prof. S. Lazebnik and her team from Illinois University for permitting to use their course material for lecture slides.