


# Face Recognition Using Deep Neural Networks




# One Shot Learning

---

- Learning from one ( or a few) examples to recognize a person.
- The idea is to understand the similarity between the detected object to a known object.

# Learning Similarity

- $d(\text{img1}, \text{img2})$  = degree of difference between images



# How can we measure similarity?

- We will find a function that quantifies a “distance” between every pair of elements in a set

Non-negativity:  $f(x, y) \geq 0$

Identity of Discernible:  $f(x, y) = 0 \iff x = y$

Symmetry:  $f(x, y) = f(y, x)$

Triangle Inequality:  $f(x, z) \leq f(x, y) + f(y, z)$

# Distance Selection \ Learning

- **Pre-defined Metrics**

Metrics which are fully specified without the knowledge of data.

E.g. Euclidian Distance:

$$f(\vec{x}, \vec{y}) = \sqrt{\sum_i (x_i - y_i)^2}$$

# Distance Selection \ Learning

- **Learned Metrics**

Metrics which can only be defined with the **knowledge** of the **data**:

- Un-Supervised Learning

Or

- Supervised Learning

# Un-supervised distance metric: :


- Example - Mahalanobis Distance :

- $f(x, y) = (x - y)^T S^{-1} (x - y)$

where  $S$  is the mean-subtracted covariance matrix of all data points.

# Un-supervised distance metric:

- 2-step procedure:
  - Apply some **supervised** domain transform:



- Then use one of the un-supervised metrics for performing the mapping.



# The One-Shot learning challenge


- There are lot of categories
  - The Number of categories is not always known
  - The number of samples in each category is small
- 
- One shot learning is relevant in the field of **computer vision**: recognize objects in images from a single example.

# Face Recognition challenges

## Verification



## Recognition



## Clustering








# Using CNN for Face Recognition

- The Idea is **to learn a function** that maps input patterns to target space.
- **non-linear** mapping that can map any input vector to its corresponding low-dimensional version.
- The distance in the target space approximates the “semantic” distance in the input space.
- Extract information about the problem from the available data, without requiring specific information about the categories.
- The training will be on pairs of samples.

# Siamese network








$x^1$



$f(x^1)$

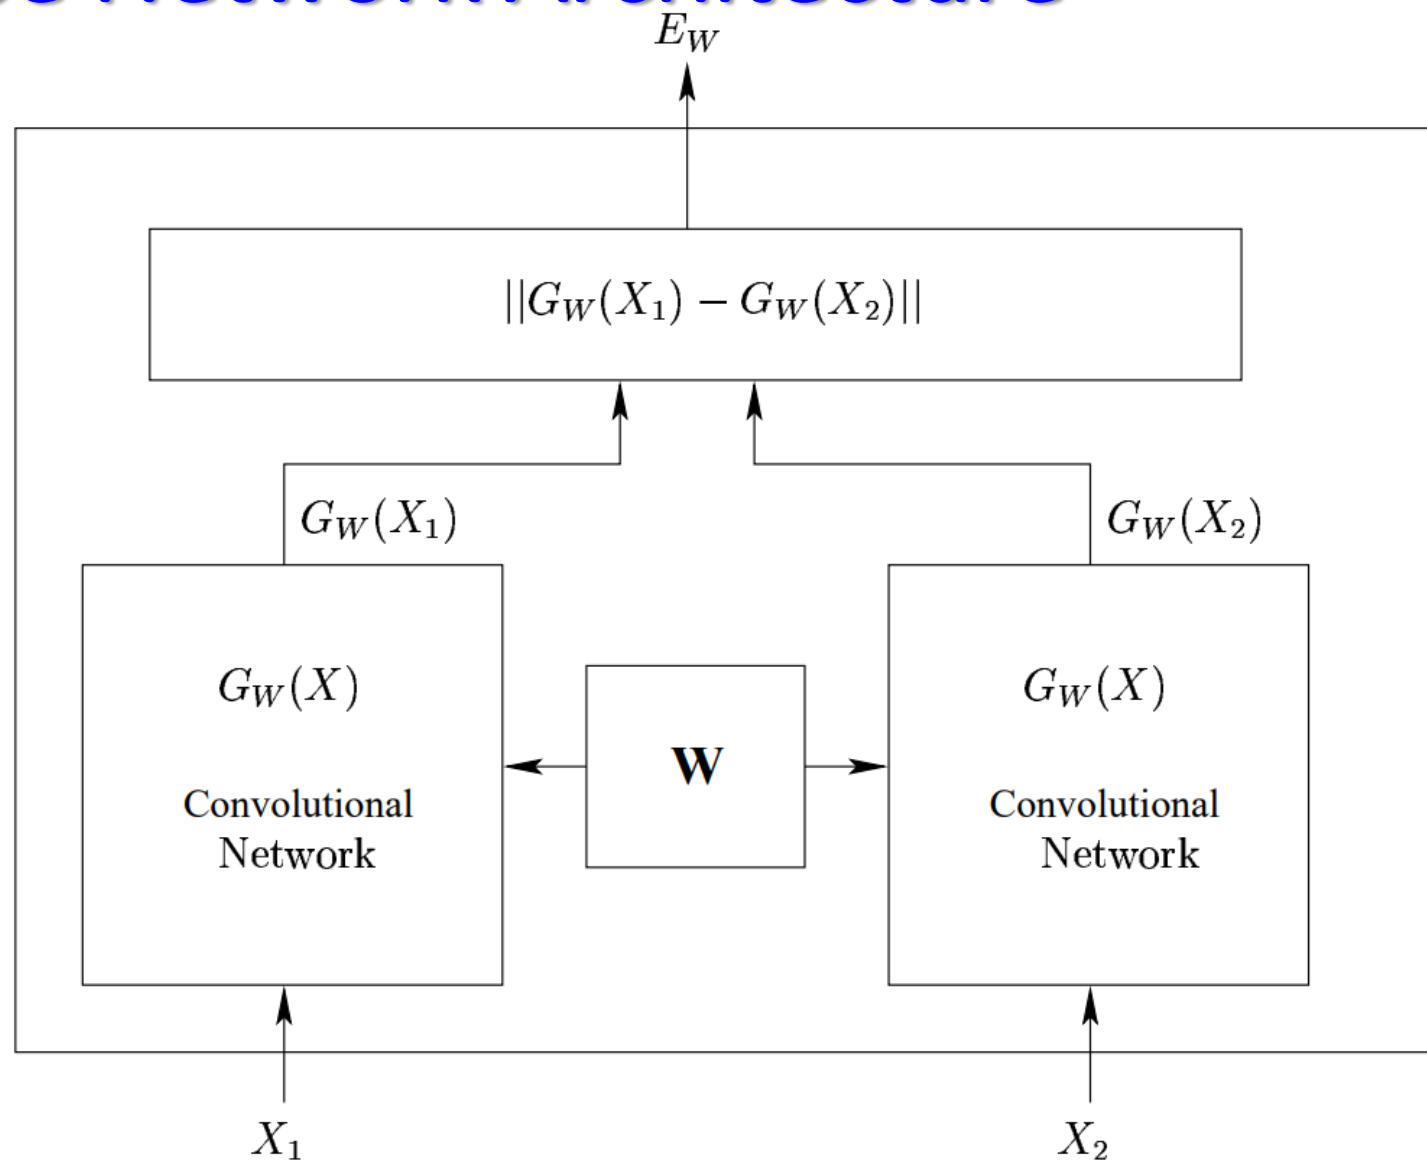


$x^2$



$f(x^2)$

# Siamese Network Architecture



# Similarity Metric

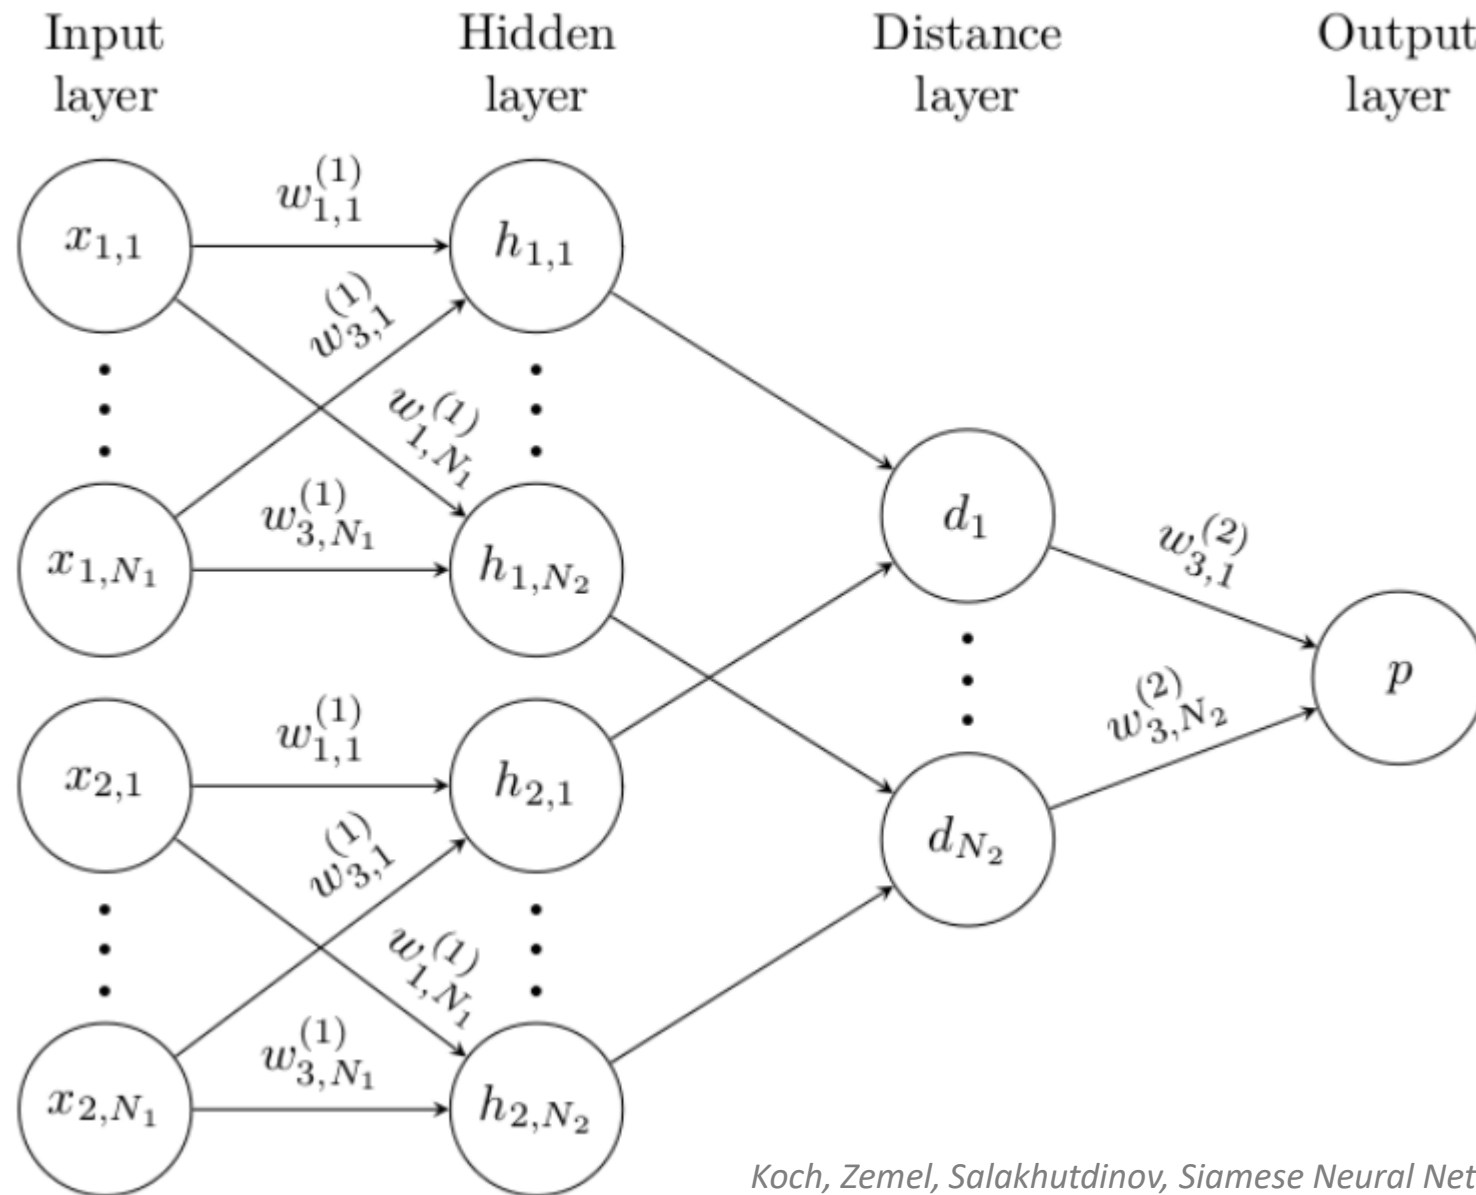
$$E_W(X_1, X_2) = ||G_W(X_1) - G_W(X_2)||$$

We seek to find a value of the parameter  $W$  such that:

$X_1, X_2$ - Same Category	$X_1, X_2$ - Different Categories
Genuine Pair	Impostor Pair
Minimizes $E_W(X_1, X_2)$	Maximizes $E_W(X_1, X_2)$

- Contrastive term is needed to ensure: not only that the energy for a pair of inputs from the same category is low, but also that the energy for a pair from different categories is large.

# Siamese Neural Networks



# Contrastive Loss Function

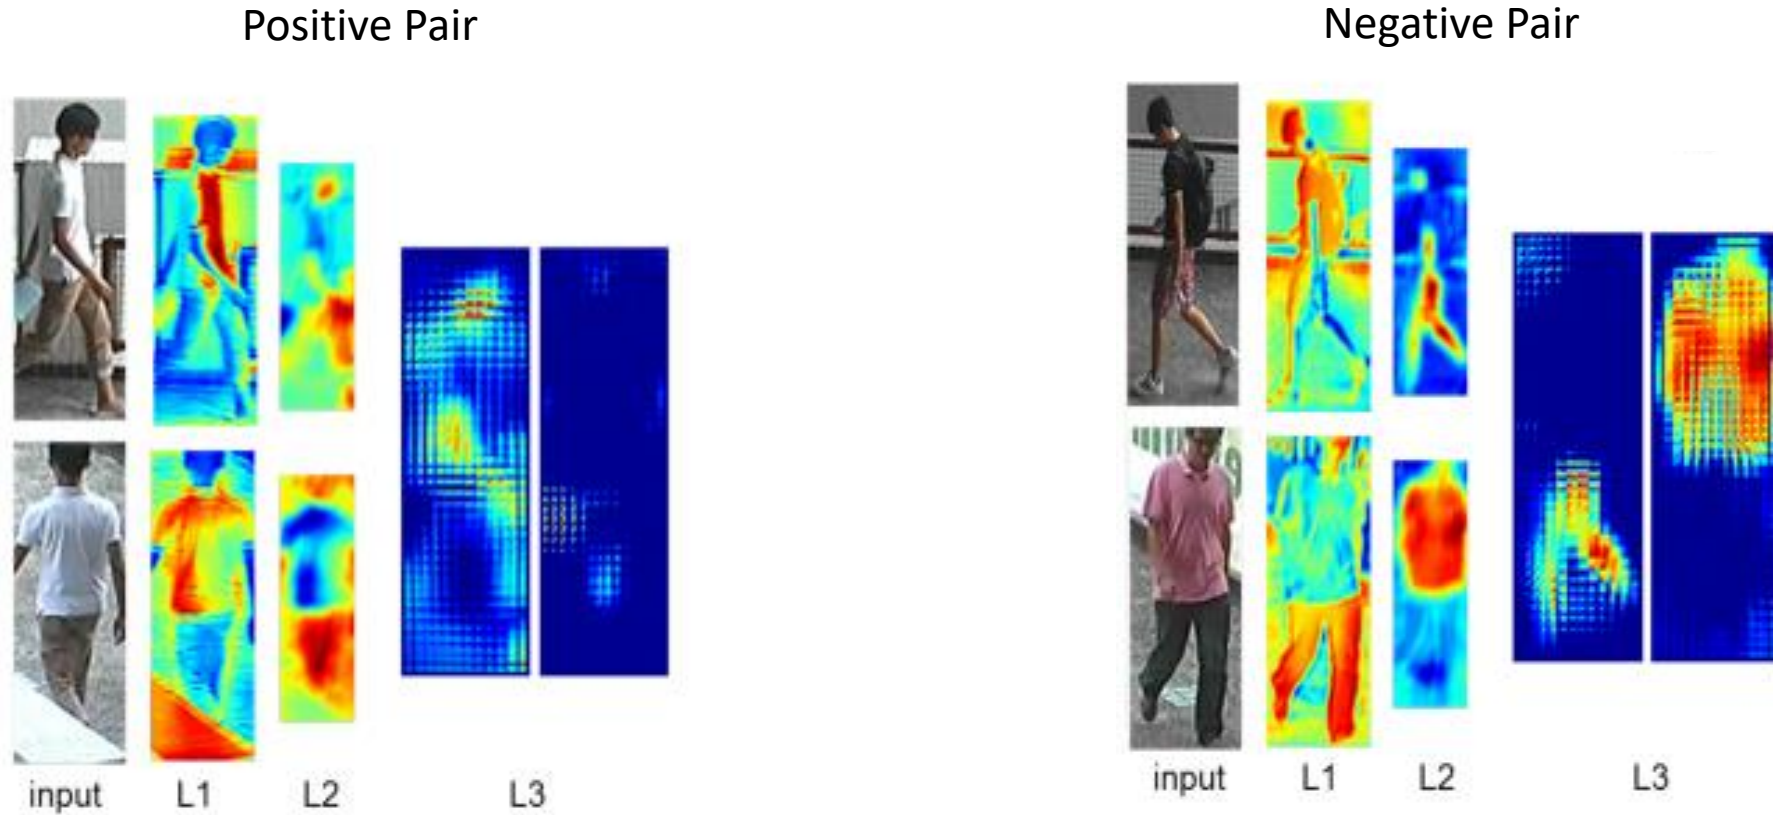
- Contrastive loss between a pair of samples  $X^i, X^j$  is defined as

$$L(X^i, X^j) = Y^{i,j} D(X^i, X^j) + (1 - Y^{i,j}) \max(0, m - D(X^i, X^j))$$

- Here  $D(X^i, X^j)$  is the Square of the Euclidean distance and  $m$  is the margin for the distance that we set for dissimilar pair of samples.
- For a batch of input samples, we take it as the average contrastive loss.



# Visualization of Learned Features



# DeepFace (Facebook, 2014)

- The conventional pipeline:

Detect  $\Rightarrow$  align  $\Rightarrow$  represent  $\Rightarrow$  classify

- **Face alignment:** Transform a face to be in a canonical pose
- **Face representation:** Find a representation of a face which is suitable for follow-up tasks (small size, computationally cheap to compare, invariant to irrelevant changes)
- 3D face modeling
- A nine-layer deep neural network
- More than 120 million parameters


# Alignment Process




(a)




(b)




(c)



(d)



(e)



(f)



(g)



(h)

# Alignment Process

- (a) 2D Alignment - detecting 6 fiducial points inside the detection crop, centered at the center of the eyes, tip of the nose and mouth locations.
- (b) 2D Alignment - aligned crop: composing the final 2D transformation.
- (c) 2D Alignment - localizing additional 67 fiducial points
- (d) 3D Alignment - The reference 3D shape transformed to the 2D-aligned crop image-plane.


# Alignment Process

(e) 3D Alignment - The 67 fiducial points induced by the 3D model that are used to direct the piece-wise affine wrapping.

(f) 2D Alignment - The final frontalized crop



# The DeepFace architecture

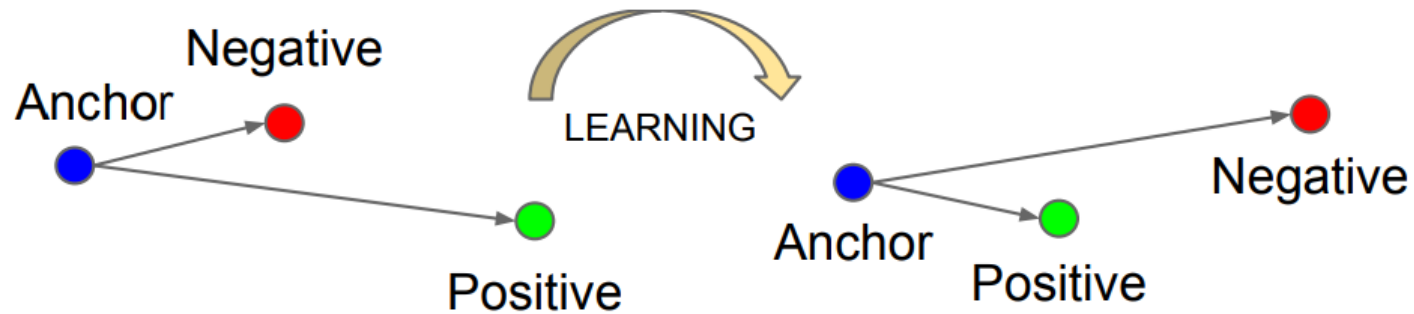


C1, M2, C3 : Extract low-level features (simple edges and texture)

L4, L5, L6: Locally connected Layers

L7, L8: Fully connected Layers

# Triplets Network



- The Triplet Loss minimizes the distance between an anchor and a positive, both of which have the same identity, and maximizes the distance between the anchor and a negative of a different identity.



# Triplet Network

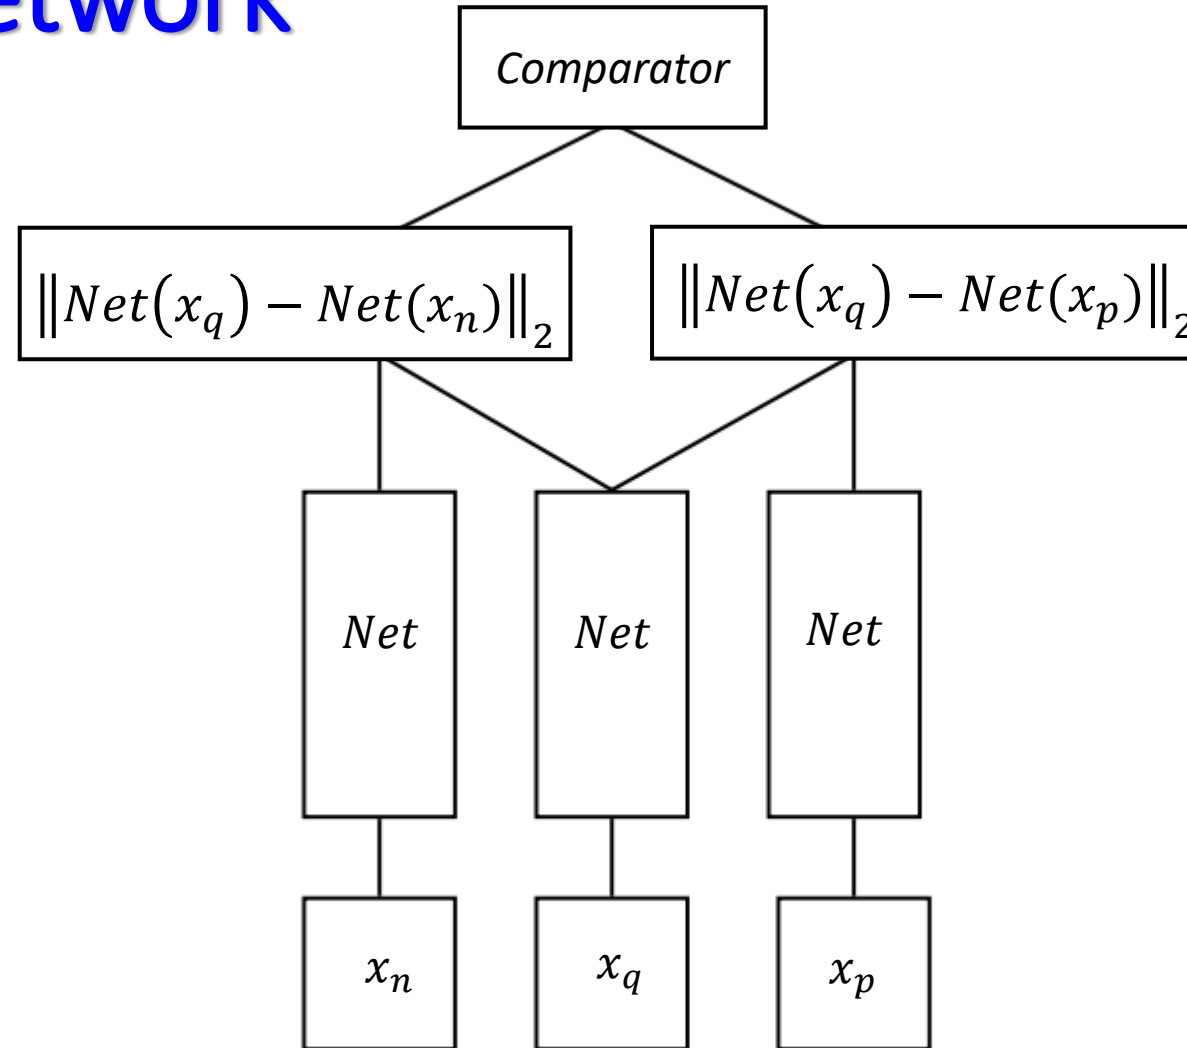


Figure 1: Triplet network structure



# Training

- We learn an embedding  $f(x)$ , from an image  $x$  into a feature space  $\mathbb{R}^d$ , such that the squared distance between all faces of the same identity is small, whereas the squared distance between a pair of face images from different identities is large.
- Loss Function:

$$L = \sum_i^N \left[ \|f(x_{q,i}) - f(x_{p,i})\|_2^2 - \|f(x_{q,i}) - f(x_{n,i})\|_2^2 + \alpha \right]_+$$

$f(\cdot)$  - the embedding function

# Triplets Selection


- We could select  $(x_{p,i})$  and  $(x_{n,i})$  :
  - (1) hard positive-  $\operatorname{argmax} (x_{p,i}) \quad \|f(x_{q,i}) - f(x_{p,i})\|_2^2$
  - (2) hard negative-  $\operatorname{argmin} (x_{q,i}) \quad \|f(x_{q,i}) - f(x_{n,i})\|_2^2$
- For fast convergence it is crucial to select triplets that violate the triplet constraint.
- It is infeasible to compute the argmin and argmax across the whole training set.
- Further, it may lead to poor training, as mislabelled and poorly imaged faces would dominate the hard positives and negatives.

# Triplet Selection

- There are two obvious choices that avoid this issue:
  - Generate triplets offline every n steps, using the most recent network checkpoint and computing the argmin and argmax on a subset of the data.
  - Generate triplets online. This can be done by selecting the hard positive/negative examples from within a mini-batch.
- To avoid local minima, we use "semi-hard" examples such that

$$\|f(x_{q,i}) - f(x_{p,i})\|_2^2 < \|f(x_{q,i}) - f(x_{n,i})\|_2^2$$

# The Model structure (after the training)



# The FaceNet architecture (NN1)

layer	size-in	size-out	kernel	param	FLPS
conv1	$220 \times 220 \times 3$	$110 \times 110 \times 64$	$7 \times 7 \times 3, 2$	9K	115M
pool1	$110 \times 110 \times 64$	$55 \times 55 \times 64$	$3 \times 3 \times 64, 2$	0	
rnorm1	$55 \times 55 \times 64$	$55 \times 55 \times 64$		0	
conv2a	$55 \times 55 \times 64$	$55 \times 55 \times 64$	$1 \times 1 \times 64, 1$	4K	13M
conv2	$55 \times 55 \times 64$	$55 \times 55 \times 192$	$3 \times 3 \times 64, 1$	111K	335M
rnorm2	$55 \times 55 \times 192$	$55 \times 55 \times 192$		0	
pool2	$55 \times 55 \times 192$	$28 \times 28 \times 192$	$3 \times 3 \times 192, 2$	0	
conv3a	$28 \times 28 \times 192$	$28 \times 28 \times 192$	$1 \times 1 \times 192, 1$	37K	29M
conv3	$28 \times 28 \times 192$	$28 \times 28 \times 384$	$3 \times 3 \times 192, 1$	664K	521M
pool3	$28 \times 28 \times 384$	$14 \times 14 \times 384$	$3 \times 3 \times 384, 2$	0	
conv4a	$14 \times 14 \times 384$	$14 \times 14 \times 384$	$1 \times 1 \times 384, 1$	148K	29M
conv4	$14 \times 14 \times 384$	$14 \times 14 \times 256$	$3 \times 3 \times 384, 1$	885K	173M
conv5a	$14 \times 14 \times 256$	$14 \times 14 \times 256$	$1 \times 1 \times 256, 1$	66K	13M
conv5	$14 \times 14 \times 256$	$14 \times 14 \times 256$	$3 \times 3 \times 256, 1$	590K	116M
conv6a	$14 \times 14 \times 256$	$14 \times 14 \times 256$	$1 \times 1 \times 256, 1$	66K	13M
conv6	$14 \times 14 \times 256$	$14 \times 14 \times 256$	$3 \times 3 \times 256, 1$	590K	116M
pool4	$14 \times 14 \times 256$	$7 \times 7 \times 256$	$3 \times 3 \times 256, 2$	0	
concat	$7 \times 7 \times 256$	$7 \times 7 \times 256$		0	
fc1	$7 \times 7 \times 256$	$1 \times 32 \times 128$	maxout p=2	103M	103M
fc2	$1 \times 32 \times 128$	$1 \times 32 \times 128$	maxout p=2	34M	34M
fc7128	$1 \times 32 \times 128$	$1 \times 1 \times 128$		524K	0.5M
L2	$1 \times 1 \times 128$	$1 \times 1 \times 128$		0	
total				140M	1.6B

# The FaceNet architecture (NN2)

type	output size	depth	#1×1	#3×3 reduce	#3×3	#5×5 reduce	#5×5	pool proj (p)	params	FLOPS
conv1 (7×7×3, 2)	112×112×64	1							9K	119M
max pool + norm	56×56×64	0						m 3×3, 2		
inception (2)	56×56×192	2		64	192				115K	360M
norm + max pool	28×28×192	0						m 3×3, 2		
inception (3a)	28×28×256	2	64	96	128	16	32	m, 32p	164K	128M
inception (3b)	28×28×320	2	64	96	128	32	64	$L_2$ , 64p	228K	179M
inception (3c)	14×14×640	2	0	128	256,2	32	64,2	m 3×3,2	398K	108M
inception (4a)	14×14×640	2	256	96	192	32	64	$L_2$ , 128p	545K	107M
inception (4b)	14×14×640	2	224	112	224	32	64	$L_2$ , 128p	595K	117M
inception (4c)	14×14×640	2	192	128	256	32	64	$L_2$ , 128p	654K	128M
inception (4d)	14×14×640	2	160	144	288	32	64	$L_2$ , 128p	722K	142M
inception (4e)	7×7×1024	2	0	160	256,2	64	128,2	m 3×3,2	717K	56M
inception (5a)	7×7×1024	2	384	192	384	48	128	$L_2$ , 128p	1.6M	78M
inception (5b)	7×7×1024	2	384	192	384	48	128	m, 128p	1.6M	78M
avg pool	1×1×1024	0								
fully conn	1×1×128	1							131K	0.1M
L2 normalization	1×1×128	0								
total									7.5M	1.6B

# The results

- Performance on Youtube Faces DB: 95.12% accuracy
- Performance on Labeled Faces in the Wild DB: 99.63% accuracy


- Sensitivity to Image Quality:

(The CNN was trained on 220x220 input images)

# Pixels	Val-Rate
1,600	37.8%
6,400	79.5%
14,400	84.5%
25,600	85.7%
65,539	86.4%



# FaceNet – Image clustering





# References

- Chopra, Hadsell, LeCun, [Learning a Similarity Metric Discriminatively, with Application to Face Verification](#), 2005
- Taigman; Yang; Ranzato, Wolf, [DeepFace: Closing the Gap to Human-Level Performance in Face Verification](#), 2014
- Schroff, Kalenichenko, Philbin, [FaceNet: A Unified Embedding for Face Recognition and Clustering](#), 2015
- Koch, Zemel, Salakhutdinov, [Siamese Neural Networks for One-shot Image Recognition](#), 2015
- Hermans, Beyer, Leibe, [In Defense of the Triplet Loss for Person Re-Identification](#), 2017
- LeCun, Chopra, Hadsell, Ranzato, [Huang, A Tutorial on Energy-Based Learning](#), 2006
  - Bell, Bala, [Learning visual similarity for product design with convolutional neural networks](#), 2015
  - Ahmed, Jones, Marks, [An improved deep learning architecture for person re-identification](#), 2015
  - Nando de Freitas, [Max-margin learning, transfer and memory networks](#), 2015
  - Jagannatha Rao, Wang, Cottrell, [A Deep Siamese Neural Network Learns the Human-Perceived Similarity Structure of Facial Expressions Without Explicit Categories](#), 2011
  - Hoffer, Ailon, [DEEP METRIC LEARNING USING TRIPLET NETWORK](#), 2015
  - Bellet, Habrard, Sebban, [A survey on metric learning for feature vectors and structured data](#), 2013
  - Andre Ng, Lecture slides on Face Recognition, Coursera : Convolution Neural Network, 2017.